

Search I:

Tree, Graph search

Hwanjo Yu


POSTECH

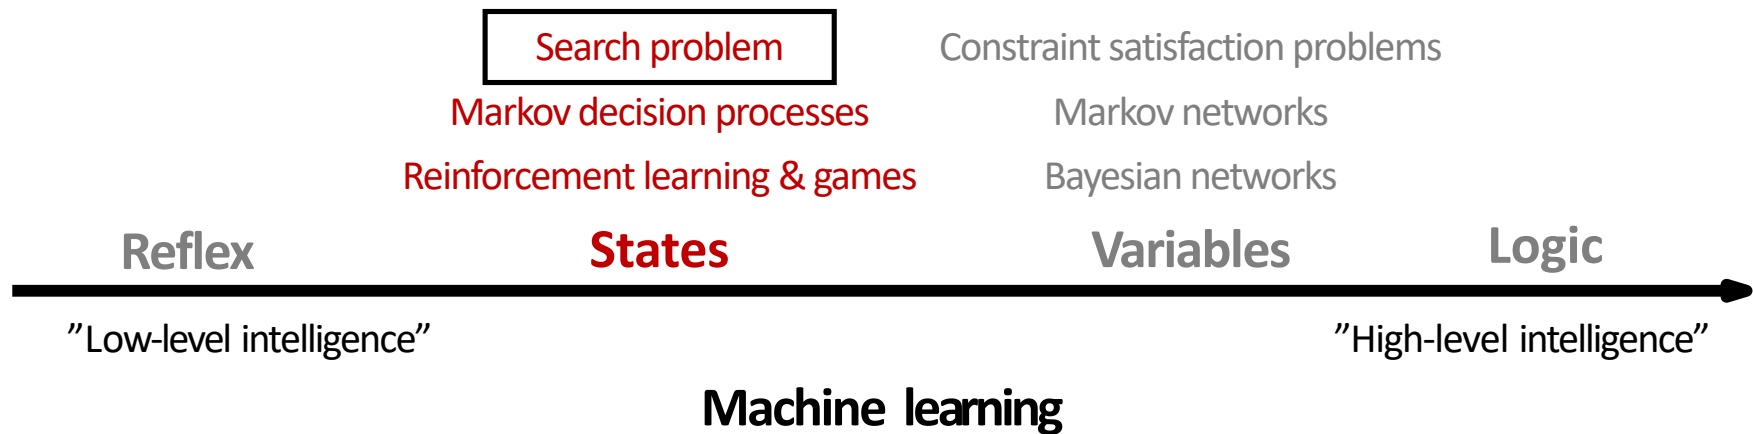
<http://di.postech.ac.kr/hwanjoyu>

Question


A farmer wants to get his cabbage, goat, and wolf across a river. He has a boat that only holds two. He cannot leave the cabbage and goat alone or the goat and wolf alone. How many river crossings does he need?

- a. 4
- b. 5
- c. 6
- d. 7
- e. No solution





Application: route finding





- **Objective:** shortest? fastest? most scenic?
- **Action:** go straight, turn left, turn right

Application: robot motion planning



- **Objective:** fastest? most energy efficient? safest?
- **Action:** translate and rotate joints

Application: solving puzzles



- **Objective:** reach a certain configuration
- **Action:** move pieces (e.g., Move12Down)

Beyond reflex

Classifier (reflex-based models):



Search problem (state-based models):



Key: models future consequences of an action!

Roadmap

Tree search

Dynamic programming

Uniform cost search




Farmer Cabbage Goat Wolf

Actions:


- F▷ F◁
- FC▷ FC◁
- FG▷ FG◁
- FW▷ FW◁

Approach: build a search tree (“what if?”)

Farmer Cabbage Goat Wolf



Search problem



Definition: search problem

- s_{start} : starting state
- $\text{Actions}(s)$: possible actions
- $\text{Cost}(s, a)$: action cost
- $\text{Succ}(s, a)$: successor
- $\text{IsEnd}(s)$: reached end state?

Transportation example


Example: transportation

- Street with blocks numbered 1 to n .
- Walking from s to $s + 1$ takes 1 minute.
- Taking a magic tram from s to $2s$ takes 2 minutes
- How to travel from 1 to n in the least time?

Backtracking search

Example: transportation

- Street with blocks numbered 1 to n .
- Walking from s to $s + 1$ takes 1 minute.
- Taking a magic tram from s to $2s$ takes 2 minutes
- How to travel from 1 to n in the least time?



If $b(= 2)$ actions per state, maximum depth is $D(= n)$ actions:

- Memory: $O(D)$ (small)
- Time: $O(b^D)$ (huge) [$2^{50} = 1,125,899,906,842,624$]

Backtracking search

Algorithm: backtracking search

- def backtrackingSearch(s , path):
 - If IsEnd(s): update minimum cost path
 - For each action $a \in \text{Actions}(s)$:
 - Extend path with Succ(s, a) and Cost(s, a)
 - Call backtrackingSearch(Succ(s, a), path)
 - Return minimum cost path
- Guarantee to find the minimum path

Depth-first search

Idea: Backtracking search + stop when find the first end state.

Assumption: zero action costs (to guarantee to find the minimum path)

- Assume action costs $\text{Cost}(s, a) = 0$

If b actions per state, maximum depth is D actions:

- Space: still $O(D)$
- Time: still $O(b^D)$ worst case, but could be much better if solutions are easy to find

Breadth-first search

Idea: explore all nodes in order of increasing depth.

Assumption: constant (same non-negative) action costs (to guarantee to find the minimum path)

- Assume action costs $\text{Cost}(s, a) = c$ for some $c \geq 0$

Legend: b actions per state, solution has d actions

- Space: now $O(b^d)$ (a lot worse!)
- Time: $O(b^d)$ (better, depends on d , not D)

DFS with iterative deepening

Idea

- Modify DFS to stop at a maximum depth
- Call DFS for maximum depths 1, 2, ...
DFS on d asks: is there a solution with d actions?

Assumption: constant (same non-negative) action costs (to guarantee to find the minimum path)

- Assume action costs $\text{Cost}(s, a) = c$ for some $c \geq 0$

Legend: b actions per state, solution size d

- Space: $O(d)$ (saved!)
- Time: $O(b^d)$ (same as BFS)

Tree search algorithms

Legend: b actions / state, solution depth d , maximum depth D

Algorithm	Action costs	Space	Time
DFS	zero	$O(D)$	$O(b^D)$
BFS	constant ≥ 0	$O(b^d)$	$O(b^d)$
DFS-ID	constant ≥ 0	$O(d)$	$O(b^d)$
Backtracking	any	$O(D)$	$O(b^D)$

- Always exponential time
- Avoid exponential space with DFS-ID


Roadmap

Tree search

Dynamic programming

Uniform cost search

Dynamic programming




Minimum cost path from state s to an end state:

$$\text{FutureCost}(s) = \begin{cases} 0 & \text{if IsEnd}(s) \\ \min_{a \in \text{Actions}(s)} [\text{Cost}(s, a) + \text{FutureCost}(\text{Succ}(s, a))] & \text{otherwise} \end{cases}$$

Motivating task

Example: route finding


- Find the minimum cost path from city 1 to city n , only moving forward. It costs c_{ij} to go from i to j



Observation: future costs only depend on current city

Dynamic programming

State: ~~past sequence of actions~~ current city




Exponential saving in time and space!

Dynamic programming

Algorithm: dynamic programming

- def DynamicProgramming(s , path):
 - If already computed for s , return cached answer
 - If IsEnd(s): return solution
 - For each action $a \in \text{Actions}(s)$:
 - Extend path with $\text{Succ}(s, a)$ and $\text{Cost}(s, a)$
 - Call $\text{DynamicProgramming}(\text{Succ}(s, a), \text{path})$
 - Return minimum cost path



Assumption: acyclicity

- The state graph defined by $\text{Action}(s)$ and $\text{Succ}(s, a)$ is acyclic.

Dynamic programming

Key idea: state

- A **state** is a summary of all the past actions sufficient to choose future actions **optimally**.

Past actions (all cities) 1 3 4 6

State (current city) 1 3 4 6

Handling additional constraints

Example: route finding

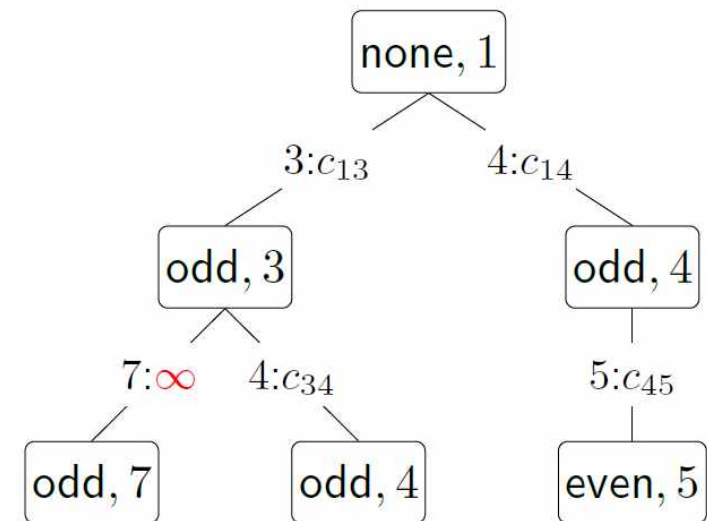
- Find the minimum cost path from city 1 to city n , only moving forward. It costs c_{ij} to go from i to j
- Constraint: Can't visit three odd cities in a row.**

State: (previous city, current city)

vs.

State: (whether previous city was odd, current city)

How many states?



Question

- Objective: travel from city 1 to city n , visiting at least 3 odd cities. What is the minimal state?

State: (# odd cities visited, current city)

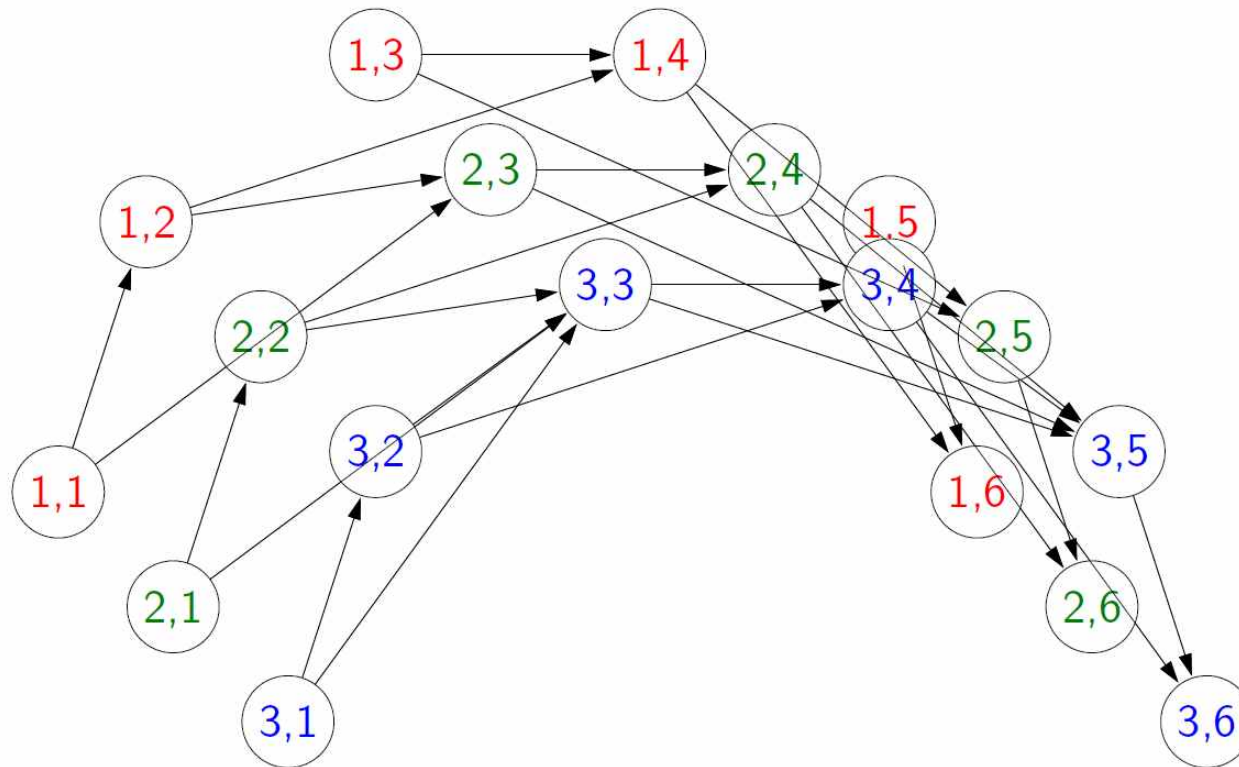
vs.

State: (min(# odd cities visited, 3), current city)

How many states?

State graph

State: (min(# odd cities visited, 3), current city)



Question

- Objective: travel from city 1 to city n , visiting more odd than even cities. What is the minimal state?

Summary

- **State:** summary of past actions sufficient (and minimal?) to choose future actions optimally
- Dynamic programming: backtracking search with memoization – potentially exponential savings

Dynamic programming only works for acyclic graphs... what if there are cycles?

Roadmap

Tree search


Dynamic programming

Uniform cost search


Uniform cost search (UCS)

Observation: prefixes of optimal path are optimal

- PastCost $PC(S)$: minimal cost from S_{start} to S



- $PC(S) = PC(S_i) + Cost(a, c)$




Key idea: state ordering

- UCS enumerates states in order of increasing past cost

Uniform cost search (UCS)

Key idea: state ordering

- UCS enumerates states in order of increasing past cost



Assume $PC(S_b) + c_3$ is minimum of

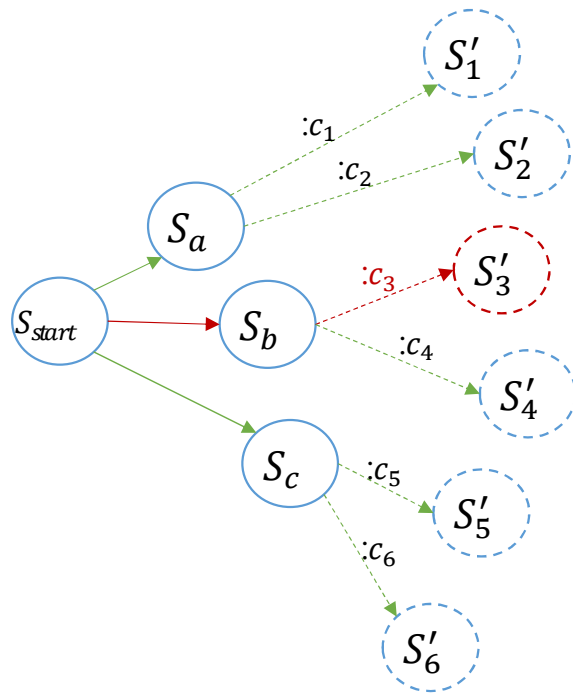
$$\begin{cases} PC(S_a) + c_1 \\ PC(S_a) + c_2 \\ PC(S_b) + c_3 \\ PC(S_b) + c_4 \\ PC(S_c) + c_5 \\ PC(S_c) + c_6 \end{cases}$$

Then, $PC(S'_3) = PC(S_b) + c_3$?

Uniform cost search (UCS)

Key idea: state ordering

- UCS enumerates states in order of increasing past cost



Assume $PC(S_b) + c_3$ is minimum of


$$\begin{cases} PC(S_a) + c_1 \\ PC(S_a) + c_2 \\ PC(S_b) + c_3 \\ PC(S_b) + c_4 \\ PC(S_c) + c_5 \\ PC(S_c) + c_6 \end{cases}$$

Then, $PC(S'_3) = PC(S_b) + c_3$?

Assumption: non-negativity


All action costs are **non-negative**: $\text{Cost}(s, a) \geq 0$

UCS: High-level strategy



- *Explored*: states we've found the optimal path to
- *Frontier (Priority Queue)*: states we've seen, still figuring out how to get there cheaply
- *Unexplored*: states we haven't seen

UCS example



Start state: A, end state: D

Minimum cost path:

A -> B -> C -> D with cost 3

Uniform cost search (UCS)

Algorithm: uniform cost search [Dijkstra, 1956]


- Add s_{start} to **frontier** (priority queue)
- Repeat until frontier is empty:
 - Remove s with smallest priority p from frontier
 - If $\text{IsEnd}(s)$: return solution
 - Add s to **explored**
 - For each action $a \in \text{Actions}(s)$:
 - Get successor $s' \leftarrow \text{Succ}(s, a)$
 - If s' already in explored: continue
 - Update **frontier** with s' and priority $p + \text{cost}(s, a)$
 - If updated: $\text{backpointers}[s'] = (s, a)$

Analysis of uniform cost search

Theorem: correctness

- When a state s is popped from the frontier and moved to explored, its priority is $\text{PastCost}(s)$, the minimum cost to s .

Proof:



DP versus UCS

N total states, n of which are closer than end state

Algorithm	Cycles?	Action costs	Time/space
DP	no	any	$O(N)$
UCS	yes	≥ 0	$O(n \log n)$

Note:

- UCS potentially explores fewer states, but requires more overhead to maintain the priority queue
- Assume number of actions per state is constant (independent of n and N)

Summary

- Tree search: memory efficient, suitable for huge state spaces but exponential worst-case running time
- State: summary of past actions sufficient to choose future actions optimally
- Graph search: dynamic programming and uniform cost search construct optimal paths (exponential savings!)
- Next time: learning action costs, searching faster with A*