

# Search II:

## Learning parameters, $A^*$

Hwanjo Yu

POSTECH

<http://di.postech.ac.kr/hwanjoyu>

# Review

Definition: search problem

- $s_{\text{start}}$  : starting state
- $\text{Actions}(s)$ : possible actions
- $\text{Cost}(s, a)$ : action cost
- $\text{Succ}(s, a)$ : successor
- $\text{IsEnd}(s)$ : reached end state?

Objective: find the minimum cost path from  $s_{\text{start}}$  to an  $s$  satisfying  $\text{IsEnd}(s)$ .

# Search

## Transportation example

- Start state: 1
- Walk action: from  $s$  to  $s + 1$  (cost: 1)
- Tram action: from  $s$  to  $2s$  (cost: 2)
- End state:  $n$



search algorithm

walk walk tram tram tram walk tram tram  
(minimum cost path)

# Learning

## Transportation example

- Start state: 1
- Walk action: from  $s$  to  $s + 1$  (cost: ?)
- Tram action: from  $s$  to  $2s$  (cost: ?)
- End state:  $n$

walk walk tram tram tram walk tram tram



learning algorithm

walk cost: 1, tram cost: 2

# Learning as an inverse problem

Forward problem (search):


$$\text{Cost}(s, a) \rightarrow (a_1, \dots, a_k)$$

Inverse problem (learning):

$$(a_1, \dots, a_k) \rightarrow \text{Cost}(s, a)$$

# Prediction (inference) problem

Input  $x$ : search problem without costs




Output  $y$ : solution path

walk walk walk

# Tweaking costs


Costs: {walk:3, tram:2}

Minimum cost path:



Costs: {walk:1, tram:3}

Minimum cost path:




# Modeling costs (simplified)

Assume costs depend only on the action:

$$\text{Cost}(s, a) = \mathbf{w}[a]$$

Candidate output path:



Path cost:

$$\text{Cost}(y) = \mathbf{w}[a_1] + \mathbf{w}[a_2] + \mathbf{w}[a_3]$$



# Learning algorithm

## Algorithm: Structured Perceptron (simplified)


- For each action:  $\mathbf{w}[a] \leftarrow 0$
- For each iteration  $t = 1, \dots T$ :
  - For each training example  $(x, y) \in D_{\text{train}}$  :
    - Compute the minimum cost path  $y'$  given  $\mathbf{w}$
    - For each action  $a \in y$ :  $\mathbf{w}[a] \leftarrow \mathbf{w}[a] - 1$
    - For each action  $a \in y'$ :  $\mathbf{w}[a] \leftarrow \mathbf{w}[a] + 1$
- Try to decrease cost of true  $y$  (from training data)
- Try to increase cost of predicted  $y'$  (from search)
- Note that if  $y = y'$ , there is no update.
- What is the implied objective (loss function)? (Notate  $\mathbf{w}[a] \rightarrow \mathbf{w} \cdot \phi(a)$ )

# Generalization to features

Costs are parametrized by feature vector

$$\text{Cost}(s, a) = \mathbf{w} \cdot \phi(s, a)$$

Example:



Path cost:

$$\text{Cost}(y) = \mathbf{w} \cdot [1, 0, 1] + \mathbf{w} \cdot [1, 2, 0] = 2 + 1 = 3$$

# Learning algorithm

## Algorithm: Structured Perceptron

- For each action:  $\mathbf{w} \leftarrow 0$
  - For each iteration  $t = 1, \dots, T$ :
    - For each training example  $(x, y) \in D_{\text{train}}$  :
      - Compute the minimum cost path  $y'$  given  $\mathbf{w}$
      - $\mathbf{w} \leftarrow \mathbf{w} - \phi(y) + \phi(y')$
  - Try to decrease cost of true  $y$  (from training data)
  - Try to increase cost of predicted  $y'$  (from search)
1. What is the implied objective (loss function)?
  2. Can you modify it using hinge loss?
  3. How the algorithm would change if using hinge loss?

# Learning algorithm

## Algorithm: Structured Perceptron

- For each action:  $\mathbf{w} \leftarrow 0$
- For each iteration  $t = 1, \dots, T$ :
  - For each training example  $(x, y) \in D_{\text{train}}$  :
    - Compute the minimum cost path  $y'$  given  $\mathbf{w}$
    - $\mathbf{w} \leftarrow \mathbf{w} - \phi(y) + \phi(y')$
- Try to decrease cost of true  $y$  (from training data)
- Try to increase cost of predicted  $y'$  (from search)


# Roadmap

Learning costs

A\* search

Relaxation


# Can uniform cost search be improved?



- Problem: UCS orders states by cost from  $s_{start}$  to  $s$
- Goal: take into account cost from  $s$  to  $s_{end}$

# Exploring states

- UCS: explore states in order of  $\text{PastCost}(s)$



- Ideal: explore in order of  $\text{PastCost}(s) + \text{FutureCost}(s)$
- A\*: explore in order of  $\text{PastCost}(s) + h(s)$

Definition: Heuristic function

- A heuristic  $h(s)$  is any estimate of  $\text{FutureCost}(s)$

# A\* search


Algorithm: A\* search [Hart/Nilsson/Raphael, 1968]

- Run uniform cost search with **modified edge costs**:

$$\text{Cost}'(s, a) = \text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s)$$

Intuition: add a penalty for how much action  $a$  takes us away from the end state

Example:



$$\text{Cost}'(C, B) = \text{Cost}(C, B) + h(B) - h(C) = 1 + (3 - 2) = 2$$




# An example heuristic

Will any heuristic work?

No.

Counterexample:




Doesn't work because of **negative modified edge costs!**

# Consistent heuristics

## Definition: consistency

- A heuristic  $h$  is **consistent** if
  - $\text{Cost}'(s, a) = \text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s) \geq 0$
  - $h(s_{\text{end}}) = 0$

- Condition 1: needed for UCS to work (triangle inequality)



- Condition 2:  $\text{FutureCost}(s_{\text{end}}) = 0$ , so match it.


# Correctness of $A^*$

## Proposition: correctness

- If  $h$  is consistent,  $A^*$  returns the minimum cost path.

# Proof of A\* Correctness

- Consider any path  $[s_0, a_1, s_1, \dots, a_L, s_L]$



- Key identity:

$$\underbrace{\sum_{i=1}^L \text{Cost}'(s_{i-1}, a_i)}_{\text{modified path cost}} = \underbrace{\sum_{i=1}^L \text{Cost}(s_{i-1}, a_i)}_{\text{original path cost}} + \underbrace{h(s_L) - h(s_0)}_{\text{constant}}$$

- Therefore, A\* (finding the minimum cost path using modified costs) solves the original problem (even though edge costs are all different!)

# Efficiency of A\*

## Theorem: efficiency of A\*

- A\* explores all states  $s$  satisfying

$$\text{PastCost}(s) \leq \text{PastCost}(s_{\text{end}}) - h(s)$$

- Note that UCS explores states in order of past cost, so that it explores every state whose past cost is less than the past cost of the end state.


**Proof:** A\* explores all  $s$  such that

$$\text{PastCost}(s) + h(s) - h(s_{\text{start}}) \leq \text{PastCost}(s_{\text{end}}) + h(s_{\text{end}}) - h(s_{\text{start}})$$

$$\text{PastCost}(s) + h(s) \leq \text{PastCost}(s_{\text{end}})$$

Interpretation: the larger  $h(s)$ , the better

# Amount explored



- If  $h(s) = 0$ , then  $A^*$  is same as UCS.
- If  $h(s) = \text{FutureCost}(s)$ , then  $A^*$  only explores nodes on a minimum cost path.
- Usually is  $h(s)$  somewhere in between.

# Admissibility

## Definition: admissibility

- A heuristic  $h(s)$  is **admissible** if  $h(s) \leq \text{FutureCost}(s)$
- Note: A\* explores all  $s$  such that  $\text{PastCost}(s) + h(s) \leq \text{PastCost}(s_{\text{end}})$

## Theorem: consistency implies admissibility

- If a heuristic  $h(s)$  is **consistent**, then  $h(s)$  is **admissible**
- Note: a consistent heuristic  $h(s) \leq \text{Cost}(s, a) + h(s')$  and  $h(s_{\text{end}}) = 0$

## Proof:

If  $h(s)$  is **not admissible**, i.e.,  $h(s) = \text{FutureCost}(s) + \alpha$  ( $\alpha > 0$ ),

Then,

$$\begin{aligned}\text{Cost}'(s', a) &= \text{Cost}(s', a) + h(s_{\text{end}}) - h(s') \text{ for } s' \text{ where } \text{Cost}(s', a) = \text{FutureCost}(s') \\ \text{Cost}'(s', a) &= \text{FutureCost}(s') - (\text{FutureCost}(s') + \alpha) < 0\end{aligned}$$

Thus,  $h(s)$  is **not consistent**

## $A^*$ with admissible heuristic

1. If  $h$  is admissible,  $h$  is consistent?
2. If  $h$  is admissible,  $A^*$  returns the minimum cost path?



# Roadmap

Learning costs

A\* search

Relaxation

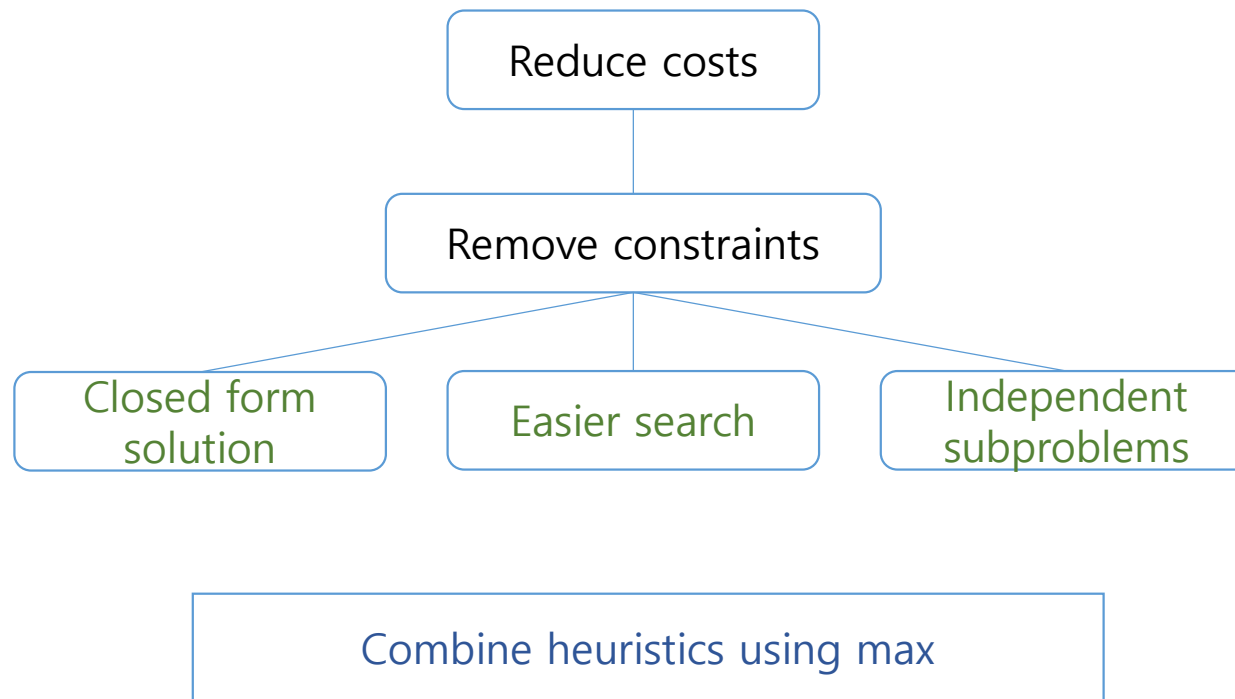
# Relaxation

Intuition: ideally, use  $h(s) = \text{FutureCost}(s)$ , but that's as hard as solving the original problem.

Key idea: relaxation

- Constraints make life hard. Get rid of them. But this is just for the heuristic!


# Relaxation overview




# Closed form solution

## Example: knock down walls

- Goal: move from triangle to circle



Hard



Easy

- Heuristic:

$$h(s) = \text{ManhattanDistance}(s, (2,5))$$

e.g.,  $h((1,1)) = 5$

# Easier search

## Example: original problem

- Start state: 1
- Walk action: from  $s$  to  $s + 1$  (cost: 1)
- Tram action: from  $s$  to  $2s$  (cost: 2)
- End state:  $n$
- Constraint: can't have more tram actions than walk actions

State: (location, #walk - #tram)

Number of states goes from  $O(n)$  to  $O(n^2)$ !

# Easier search

## Example: original problem

- Start state: 1
- Walk action: from  $s$  to  $s + 1$  (cost: 1)
- Tram action: from  $s$  to  $2s$  (cost: 2)
- End state:  $n$
- ~~Constraint: can't have more tram actions than walk actions~~

Original state: (location, ~~#walk - #tram~~)

Relaxed state: location

# Easier search


- Compute relaxed  $\text{FutureCost}'(\text{location})$  for each location  $(1, \dots, n)$  using dynamic programming or UCS
- Modify UCS to compute all past costs in reversed relaxed problem (equivalent to future costs in relaxed problem!)

## Example: reversed relaxed problem

- Start state:  $n$
  - Walk action: from  $s$  to  $s - 1$  (cost: 1)
  - Tram action: from  $s$  to  $s/2$  (cost: 2)
  - End state: 1
- 
- Define heuristic for original problem:  
$$h((\text{location}, \#walk - \#tram)) = \text{FutureCost}'(\text{location})$$

# Independent subproblems

[8 puzzle]



- Original problem: tiles **cannot** overlap (constraint)
- Relaxed problem: tiles **can** overlap (no constraint)
- Relaxed solution: 8 independent problems, each in closed form

## Key idea: independence

- Relax original problem into independent subproblems



# General framework

## Removing constraints


(knock down walls, walk/tram freely, overlap pieces)



## Reducing edge costs

(from  $\infty$  to some finite cost)

Example:



Original:  $\text{Cost}((1,1), \text{East}) = \infty$

Relaxed:  $\text{Cost}'((1,1), \text{East}) = 1$

# General framework

## Definition: relaxed search problem

- A **relaxation**  $P'$  of a search problem  $P$  has costs that satisfy:

$$\text{Cost}'(s, a) \leq \text{Cost}(s, a)$$

Note:  $\text{Cost}'(s, a)$  here is **different** from that in the definition of *consistent heuristics*: A heuristic  $h$  is **consistent** if  $h(s_{\text{end}}) = 0$  and  $\text{Cost}'(s, a) = \text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s) \geq 0$

## Definition: relaxed heuristic

- Given a relaxed search problem  $P'$ , define the **relaxed heuristic**  $h(s) = \text{FutureCost}'(s)$ , the minimum cost from  $s$  to an end state using  $\text{Cost}'(s)$ .

# General framework

## Theorem: consistency of relaxed heuristics

- Suppose  $h(s) = \text{FutureCost}'(s)$  for some relaxed problem  $P'$ . Then  $h(s)$  is a consistent heuristic.

Proof:

- $h(s) \leq \text{Cost}'(s, a) + h(\text{Succ}(s, a)) \leq \text{Cost}(s, a) + h(\text{Succ}(s, a))$

# Tradeoff

Correctness:

- Edge costs are lower => consistent heuristic
- Example: removing constraints

Efficiency:

- $h(s) = \text{FutureCost}'(s)$  must be easy to compute
- Examples: closed form, easier search, independent subproblems

# Max of two heuristics

How do we combine two heuristics?

Proposition: max heuristics

- Suppose  $h_1(s)$  and  $h_2(s)$  are consistent.
- Then  $h(s) = \max\{h_1(s), h_2(s)\}$ .

Proof: exercise

# Summary

- Structured Perceptron (reverse engineering): learn cost functions (search + learning)
- A\*: add in heuristic estimate of future costs
- Relaxation (breaking the rules): framework for producing consistent heuristics
- Next time: when actions have unknown consequences...