

# Reinforcement Learning

Hwanjo Yu


POSTECH

<http://di.postech.ac.kr/hwanjoyu>

# Review: MDPs

## Definition: Markov Decision Process

- States: the set of states
- $s_{\text{start}} \in \text{States}$  : starting state
- $\text{Actions}(s)$  : possible actions from state  $s$
- $T(s, a, s')$  : probability of  $s'$  if take action  $a$  in state  $s$
- $\text{Reward}(s, a, s')$  : reward for the transition  $(s, a, s')$
- $\text{IsEnd}(s)$  : whether  $s$  is an end of game
- $0 \leq \gamma \leq 1$  : discount factor (default: 1)



## Review: MDPs

- Following a **policy**  $\pi(s)$  produces a path (**episode**)

$$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$$

- Value** function  $V_\pi(s)$ : expected utility if follow  $\pi$  from state  $s$

$$V_\pi(s) = \begin{cases} 0 & \text{If IsEnd}(s) \\ Q_\pi(s, \pi(s)) & \text{otherwise.} \end{cases}$$

- Q-value function  $Q_\pi(s, a)$ : expected utility if first take action  $a$  from state  $s$  and then follow  $\pi$


$$Q_\pi(s, a) = \sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_\pi(s')]$$

# Unknown transitions and rewards

## Definition: Markov Decision Process

- States: the set of states
- $s_{\text{start}} \in \text{States}$  : starting state
- $\text{Actions}(s)$  : possible actions from state  $s$
- $T(s, a, s') : ?$
- $\text{Reward}(s, a, s') : ?$
- $\text{IsEnd}(s)$  : whether at end of game
- $0 \leq \gamma \leq 1$  : discount factor (default: 1)

Reinforcement learning!



# Mystery game

## Example: mystery game

- For each round  $r = 1, 2, \dots$ 
  - You choose **A** or B.
  - You move to a new state and get some rewards.
- You should take good actions to get rewards, but in order to know which actions are good, we need to explore and try different actions.

# From MDPs to reinforcement learning

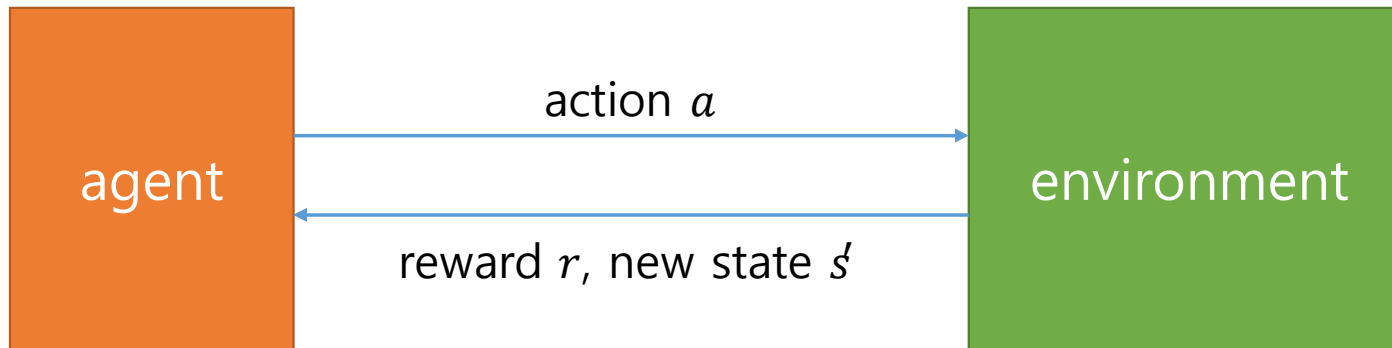
## Markov decision process (offline)

- Have mental model of how the world works.
- Find policy to collect maximum rewards.

## Reinforcement learning (online)

- Don't know how the world works.
- Perform actions in the world to find out and collect rewards.

# Reinforcement learning framework



## Algorithm: reinforcement learning template

- For  $t = 1, 2, 3, \dots$ 
  - Choose action  $a_t = \pi_{\text{act}}(s_{t-1})$  (**how?**)
  - Receive reward  $r_t$  and observe new state  $s_t$
  - Update parameters (**how?**)



# Roadmap

Monte Carlo methods

SARSA, Q-learning

Exploitation / exploration

Function approximation

# Model-based Monte Carlo

## Key idea: model-based learning

- Estimate the MDP:  $T(s, a, s')$  and  $\text{Reward}(s, a, s')$

Data:  $s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$

- Transitions:

$$\hat{T}(s, a, s') = \frac{\# \text{ times } (s, a, s') \text{ occurs}}{\# \text{ times } (s, a) \text{ occurs}}$$

- Rewards:

$$\widehat{\text{Reward}}(s, a, s') = \text{average of } r \text{ in } (s, a, r, s')$$

# Model-based Monte Carlo

## Example: model-based Monte Carlo

- Data (following policy  $\pi$ ):
  - **S1**; A, 3, **S1**; B, 0, **S1**; A, 5, **S1**; A, 7, **S2**
- Estimate:
  - $\hat{T}(S1, A, S1) = \frac{2}{3}$
  - $\hat{T}(S1, A, S2) = \frac{1}{3}$
  - $\widehat{\text{Reward}}(S1, A, S1) = \frac{1}{2}(3 + 5) = 4$
  - $\widehat{\text{Reward}}(S1, A, S2) = 7$
- Estimates converge to true values (under certain conditions)

# Problem

- Data (following policy  $\pi$ ):
  - **S1**; A, 3, **S2**; B, 0, **S1**; A, 5, **S1**; A, 7, **S1**
- Problem:
  - won't even see  $(s, a)$  if  $a \neq \pi(s)$
- Solution:
  - Need  $\pi$  to explore explicitly (more on this later)

## Key idea: exploration

- To do reinforcement learning, need to explore the state space.

## From model-based to model-free

$$\hat{Q}_{\text{opt}}(s, a) = \sum_{s'} \hat{T}(s, a, s') [\widehat{\text{Reward}}(s, a, s') + \gamma \hat{V}_{\text{opt}}(s')]$$

- All that matters for prediction is (estimate of)  $Q_{\text{opt}}(s, a)$

Key idea: model-free learning

- Try to estimate  $Q_{\text{opt}}(s, a)$  directly.

# Model-free Monte Carlo

Data (following policy  $\pi$ ):

$$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$$

Recall:

- $Q_\pi(s, a)$  is expected utility starting at  $s$ , first taking action  $a$ , and then following policy  $\pi$ .

Utility:

- $u_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$

Estimate:

- $\hat{Q}_\pi(s, a) = \text{average of } u_t \text{ where } s_{t-1} = s, a_t = a$

# Model-free Monte Carlo

## Example: model-free Monte Carlo

- Data (following policy  $\pi$ ):
  - **S1**; A, 3, **S1**; B, 0, **S1**; A, 5, **S1**; A, 7, **S2**
- Estimate (assume  $\gamma = 1$ ):
  - $\hat{Q}_{\pi}(S1, A) = \frac{1}{3}(15 + 12 + 7) \cong 11.33$
- Note: we are estimating  $Q_{\pi}$  now, not  $Q_{opt}$ ; can use policy improvement to get new policy
- Caveat: converges, but still need follow  $\pi$  that explores (**on-policy** algorithm whereas model-based Monte Carlo is **off-policy**)

# Model-free Monte Carlo (equivalences)

Data (following policy  $\pi$ ):

$$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$$

Original formulation:

- $\hat{Q}_\pi(s, a) = \text{average of } u_t \text{ where } s_{t-1} = s, a_t = a$

Equivalent formulation (convex combination):

- On each  $(s, a, u)$ :
  - $\eta = \frac{1}{1 + (\# \text{ updates to } (s, a))}$
  - $\hat{Q}_\pi(s, a) \leftarrow (1 - \eta)\hat{Q}_\pi(s, a) + \eta u$



# Model-free Monte Carlo (equivalences)

Equivalent formulation (convex combination):


- On each  $(s, a, u)$ :
  - $\hat{Q}_\pi(s, a) \leftarrow (1 - \eta)\hat{Q}_\pi(s, a) + \eta u$

Equivalent formulation (stochastic gradient):

- On each  $(s, a, u)$ :
  - $\hat{Q}_\pi(s, a) \leftarrow \hat{Q}_\pi(s, a) - \eta \left[ \underbrace{\hat{Q}_\pi(s, a)}_{\text{prediction}} - \underbrace{u}_{\text{target}} \right]$
- Implied objective: least squares regression:

$$\min_{\hat{Q}_\pi} \sum_{(s,a,u)} (\hat{Q}_\pi(s, a) - u)^2$$

# Volcanic model-free Monte Carlo



# Roadmap

Monte Carlo methods

SARSA, Q-learning

Exploitation / exploration

Function approximation

# SARSA

Data (following policy  $\pi$ ):

$$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$$

Algorithm: model-free Monte Carlo updates

- When receive  $(s, a, u)$ :
  - $\hat{Q}_\pi(s, a) \leftarrow (1 - \eta)\hat{Q}_\pi(s, a) + \eta \underbrace{u}_{\text{data}}$

Algorithm: SARSA

- When receive  $(s, a, r, s', a')$ :
  - $\hat{Q}_\pi(s, a) \leftarrow (1 - \eta)\hat{Q}_\pi(s, a) + \eta \left[ \underbrace{r}_{\text{data}} + \gamma \underbrace{\hat{Q}_\pi(s', a')}_{\text{estimate}} \right]$

# Model-free Monte Carlo vs SARSA

$$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$$

- SARSA uses estimate  $\hat{Q}_\pi(s, a)$  instead of just raw data  $u$ .
- $u$  is only based on one path, so could have large variance, need to wait until end
- $\hat{Q}_\pi(s', a')$  based on estimate, which is more stable, update immediately

# Question

Which of the following algorithms allows you to estimate  $Q_{\text{opt}}(s, a)$  (select all that apply)?

- a. model-based Monte Carlo
- b. model-free Monte Carlo
- c. SARSA

# Q-learning

**Problem:** model-free Monte Carlo and SARSA only estimate  $Q_\pi$ , but want  $Q_{\text{opt}}$  to act optimally

| <b>Output</b>    | <b>MDP</b>        | <b>reinforcement learning</b> |
|------------------|-------------------|-------------------------------|
| $Q_\pi$          | policy evaluation | model-free Monte Carlo, SARSA |
| $Q_{\text{opt}}$ | value iteration   | <b>Q-learning</b>             |

# Q-learning

- MDP recurrence (Bellman optimality equation):

$$Q_{\text{opt}}(s, a) = \sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_{\text{opt}}(s')]$$

Algorithm: Q-learning [Watkins/Dayan, 1992]

- On each  $(s, a, r, s')$ :

$$\hat{Q}_{\text{opt}}(s, a) \leftarrow (1 - \eta) \underbrace{\hat{Q}_{\text{opt}}(s, a)}_{\text{prediction}} + \eta \underbrace{[r + \gamma \hat{V}_{\text{opt}}(s')]}_{\text{target}}$$

- Recall:

$$\hat{V}_{\text{opt}}(s') = \max_{a' \in \text{Actions}(s')} \hat{Q}_{\text{opt}}(s', a')$$



# Off-policy vs On-policy

## **On-policy:**

- evaluate or improve the data-generating policy
- model-free Monte Carlo, SARSA

## **Off-policy**

- evaluate or learn using data from another policy
- Model-based Monte Carlo, Q-learning

# Reinforcement learning algorithms

| Algorithm               | Estimating             | Based on                    |
|-------------------------|------------------------|-----------------------------|
| Model-based Monte Carlo | $\hat{T}, \hat{R}$     | $s_0, a_1, r_1, s_1, \dots$ |
| Model-free Monte Carlo  | $\hat{Q}_\pi$          | $\mu$                       |
| SARSA                   | $\hat{Q}_\pi$          | $r + \hat{Q}_\pi$           |
| Q-Learning              | $\hat{Q}_{\text{opt}}$ | $r + \hat{Q}_{\text{opt}}$  |

# Roadmap

Monte Carlo methods

SARSA, Q-learning

Exploitation / exploration

Function approximation

# Exploration

## Algorithm: reinforcement learning template


- For  $t = 1, 2, 3, \dots$ 
  - Choose action  $a_t = \pi_{\text{act}}(s_{t-1})$  (how?)
  - Receive reward  $r_t$  and observe new state  $s_t$
  - Update parameters (how?)

$$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$$

- Which **exploration policy**  $\pi_{\text{act}}$  to use?

# No exploration, all exploitation


- **Attempt 1:** Set  $\pi_{\text{act}}(s) = \arg \max_{a \in \text{Actions}(s)} \hat{Q}_{\text{opt}}(s, a)$



- **Problem:**  $\hat{Q}_{\text{opt}}(s, a)$  estimates are inaccurate, **too greedy!**

# No exploitation, all exploration

- **Attemp 2:** Set  $\pi_{\text{act}}(s) = \text{random from Actions}(s)$




- **Problem:** average utility is low because exploration is **not guided**

# $\epsilon$ -greedy

Algorithm:  $\epsilon$ -greedy policy

$$\pi_{\text{act}}(s) = \begin{cases} \arg \max_{a \in \text{Actions}(s)} \hat{Q}_{\text{opt}}(s, a) & \text{probability } 1 - \epsilon \\ \text{random from Actions}(s) & \text{probability } \epsilon \end{cases}$$



# Roadmap

Monte Carlo methods

SARSA, Q-learning

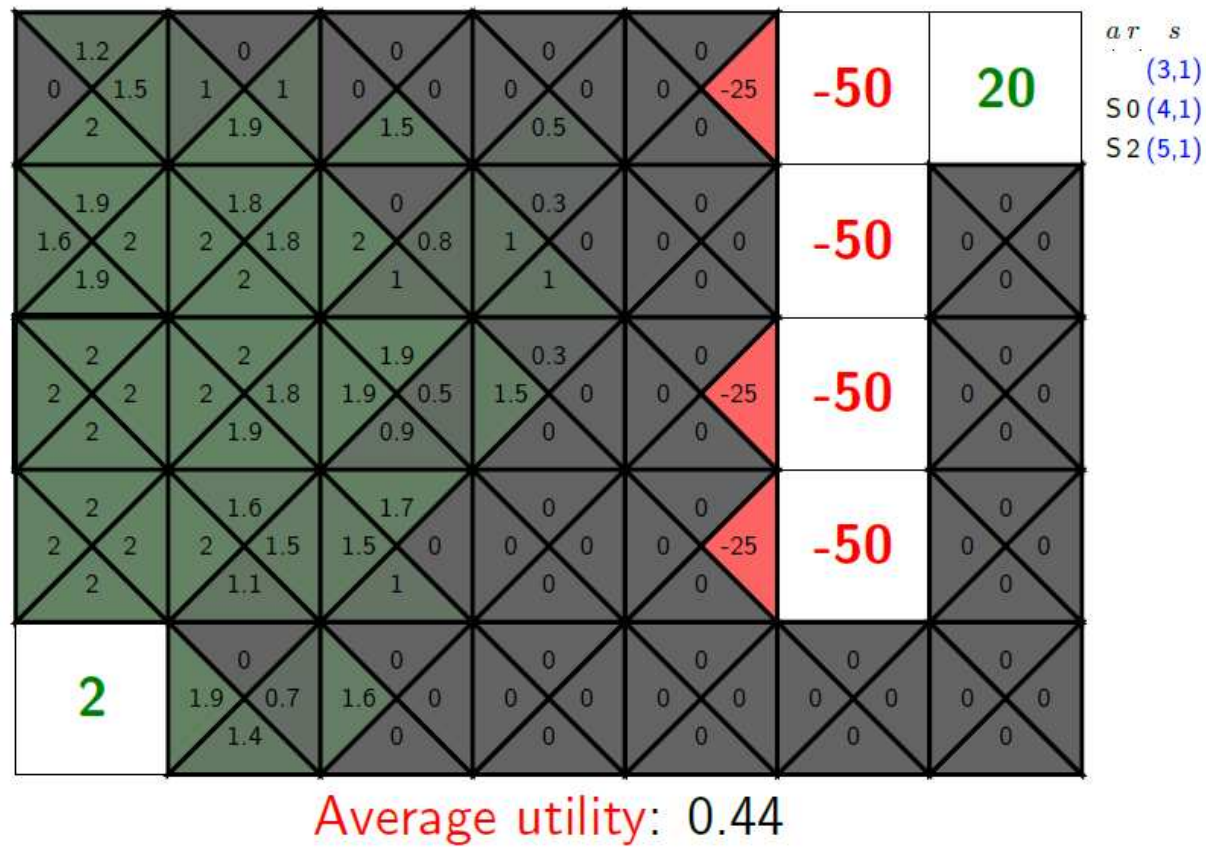
Exploitation / exploration

Function approximation



# Generalization

Problem: large state spaces, hard to explore



# Q-learning

- Stochastic gradient update:

$$\hat{Q}_{\text{opt}}(s, a) \leftarrow \hat{Q}_{\text{opt}}(s, a) - \eta \left[ \underbrace{\hat{Q}_{\text{opt}}(s, a)}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{\text{opt}}(s'))}_{\text{target}} \right]$$

- This is **rote learning**: every  $\hat{Q}_{\text{opt}}(s, a)$  has a different value
- **Problem**: doesn't generalize to unseen states/actions

# Function approximation

Key idea: linear regression model

- Define **features**  $\phi(s, a)$  and **weights**  $\mathbf{w}$ :

$$\hat{Q}_{\text{opt}}(s, a; \mathbf{w}) = \mathbf{w} \cdot \phi(s, a)$$

Example: features for volcano crossing

- $\phi_1(s, a) = \mathbf{1}[a = W]$
- $\phi_2(s, a) = \mathbf{1}[a = E]$
- ...
- $\phi_7(s, a) = \mathbf{1}[s = (5, *)]$
- $\phi_8(s, a) = \mathbf{1}[s = (*, 6)]$
- ...

# Function approximation

## Algorithm: Q-learning with function approximation

- On each  $(s, a, r, s')$ :

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \left[ \underbrace{\hat{Q}_{\text{opt}}(s, a; \mathbf{w})}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{\text{opt}}(s'))}_{\text{target}} \right] \phi(s, a)$$

Implied objective function:

$$\min_{\mathbf{w}} \sum_{(s,a,r,s')} \left( \underbrace{\hat{Q}_{\text{opt}}(s, a; \mathbf{w})}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{\text{opt}}(s'))}_{\text{target}} \right)^2$$

# Covering the unknown

- $\epsilon$ -greedy: balance the exploration/exploitation tradeoff
- Function approximation: can generalize to unseen states

# Summary: MDP and reinforcement learning

**MDP:** cope with uncertainty (unlike search problems)

- Solutions are **policies** rather than paths
- **Policy evaluation** computes policy value (expected utility)
- **Value iteration** computes optimal value (maximum expected utility) and optimal policy

**Reinforcement learning:** learn and take actions online

- **Monte Carlo:** estimate transitions, rewards, Q-values from data only
- **SARSA, Q-learning:** estimate Q-values from data and previous estimation
- **Exploitation / exploration:** balance learning and maximizing utility
- **Function approximation:** use machine learning to generalize to unseen states

# Challenges in reinforcement learning

Binary classification (sentiment classification, SVMs):

- Stateless, full feedback

Reinforcement learning (flying helicopters, Q-learning):

- Stateful, partial feedback

Key idea: partial supervision

- Reward feedback, but not given the solution directly.

Key idea: state


- Rewards depend on previous actions => can have delayed rewards.

# Crawling robot

Goal: maximize distance travelled by robot

Markov decision process (MDP):

- States: positions (4 possibilities) for each of 2 servos
- Actions: choose a servo, move it up/down
- Transitions: move into new position (unknown)
- Rewards: distance travelled (unknown)





# Deep reinforcement learning

Playing Atari [Google DeepMind, 2013]:



- Just use a neural network for  $\hat{Q}_{\text{opt}}(s, a)$
- Last 4 frames (images)  $\Rightarrow$  3-layer NN  $\Rightarrow$  keystroke
- $\epsilon$ -greedy, train over 10M frames with 1M replay memory
- <https://www.youtube.com/watch?v=V1eYniJ0Rnk>

# Deep reinforcement learning

- Policy gradient: train a policy  $\pi(a|s)$  (say, a neural network) to directly maximize expected reward

- Google DeepMind's AlphaGo (2016)



- Andrej Karpathy's blog post


<http://karpathy.github.io/2016/05/31/rl>

# Application


- Autonomous helicopters: control helicopter to do maneuvers in the air
- Backgammon: TD-Gammon plays 1-2 million games against itself, human-level performance
- Elevator scheduling; send which elevators to which floors to maximize throughput of building
- Managing datacenters; actions: bring up and shut down machine to minimize time/cost

# State-based Models

- **Search problems:** you control everything



- **MDP, RL:** against nature (e.g., Blackjack)



## Next time...

- **Adversarial games:** against opponent (e.g., chess)

