

Factor Graph and Constraint Satisfaction Problems (CSPs) 2

Hwanjo Yu

POSTECH

<http://di.postech.ac.kr/hwanjoyu>

Roadmap

Modeling

Definitions

Examples

Backtracking (exact) search

Dynamic ordering


Arc consistency

Approximate search

Beam search

Local search

Review: factor graph and CSPs



Definition: factor graph

- Variables:

$X = (X_1, \dots, X_n)$, where $X_i \in \text{Domain}_i$

- Factors:

f_1, \dots, f_m , with each $f_j(X) \geq 0$

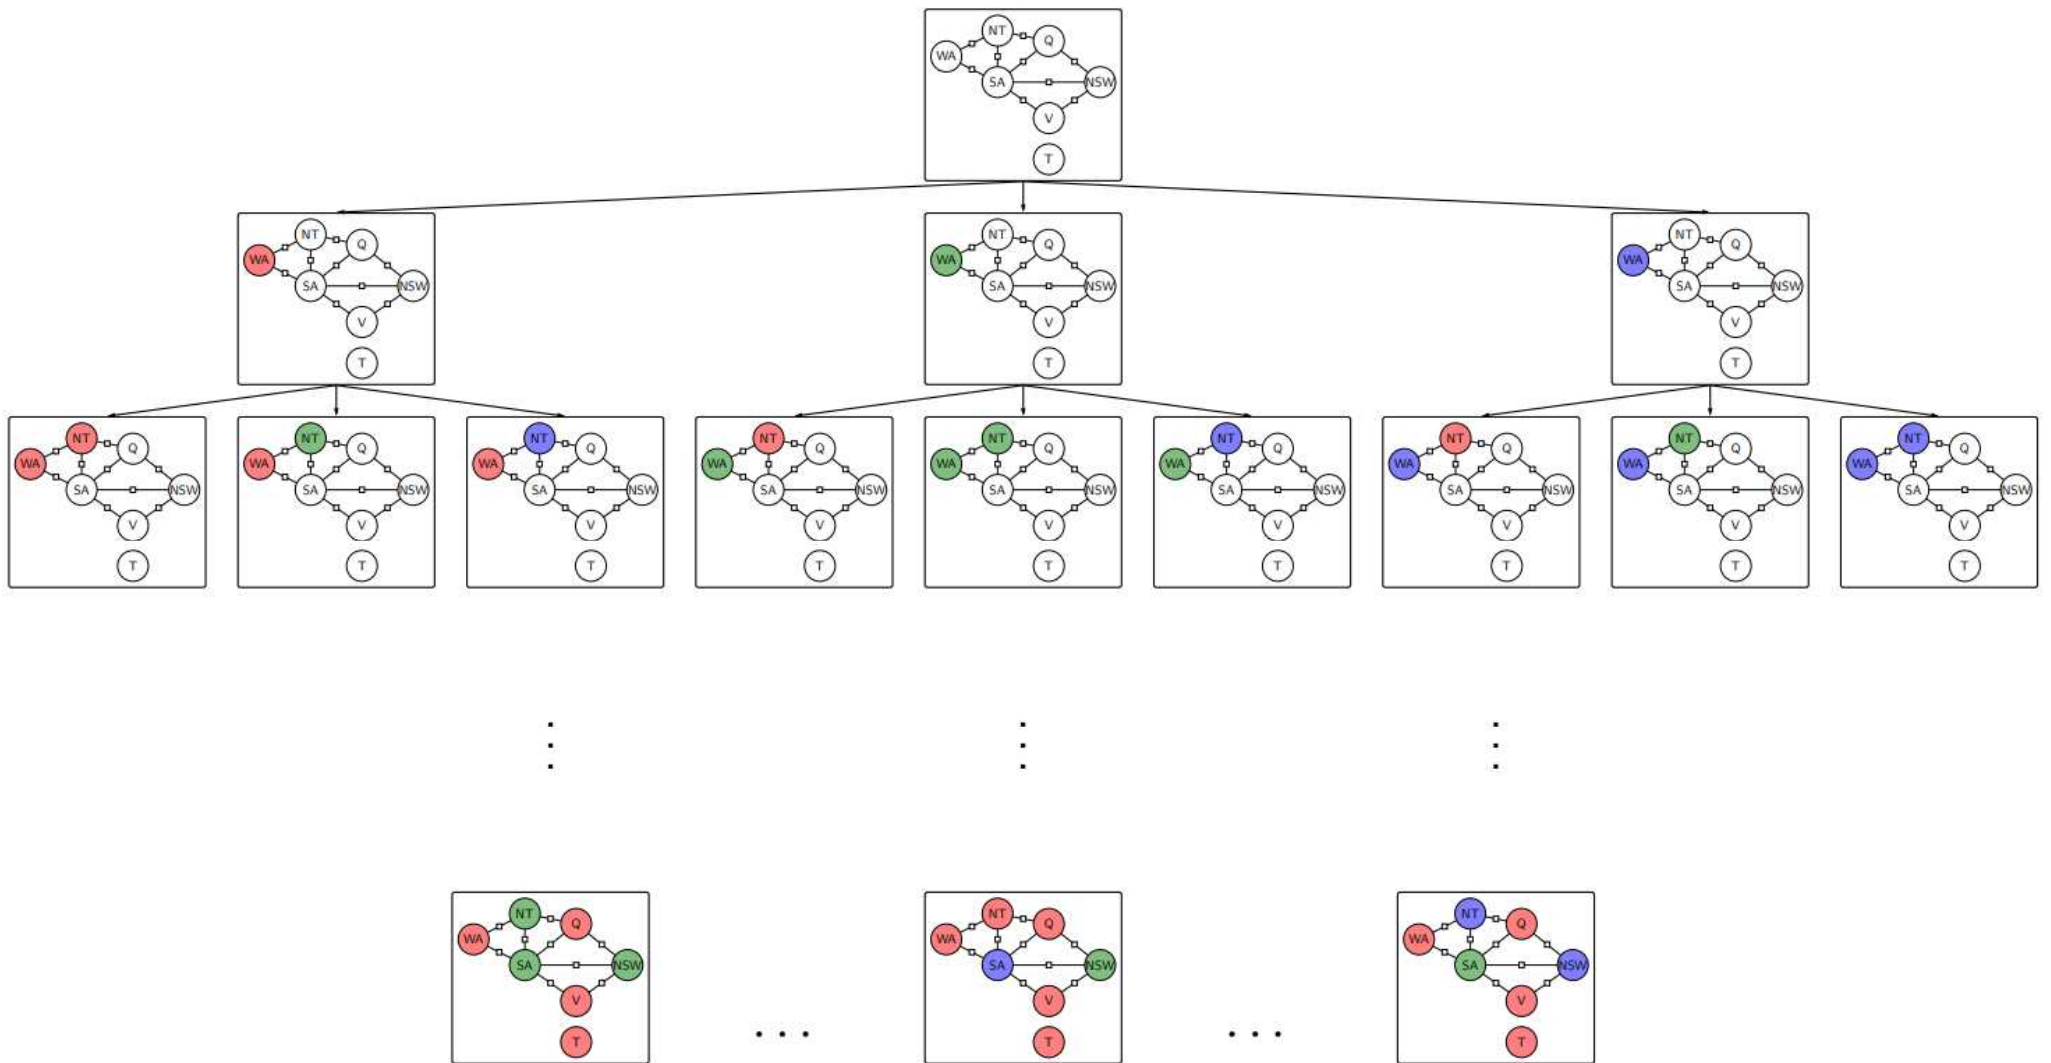
Definition: assignment weight

- Each **assignment** $x = (x_1, \dots, x_n)$ has a **weight**:

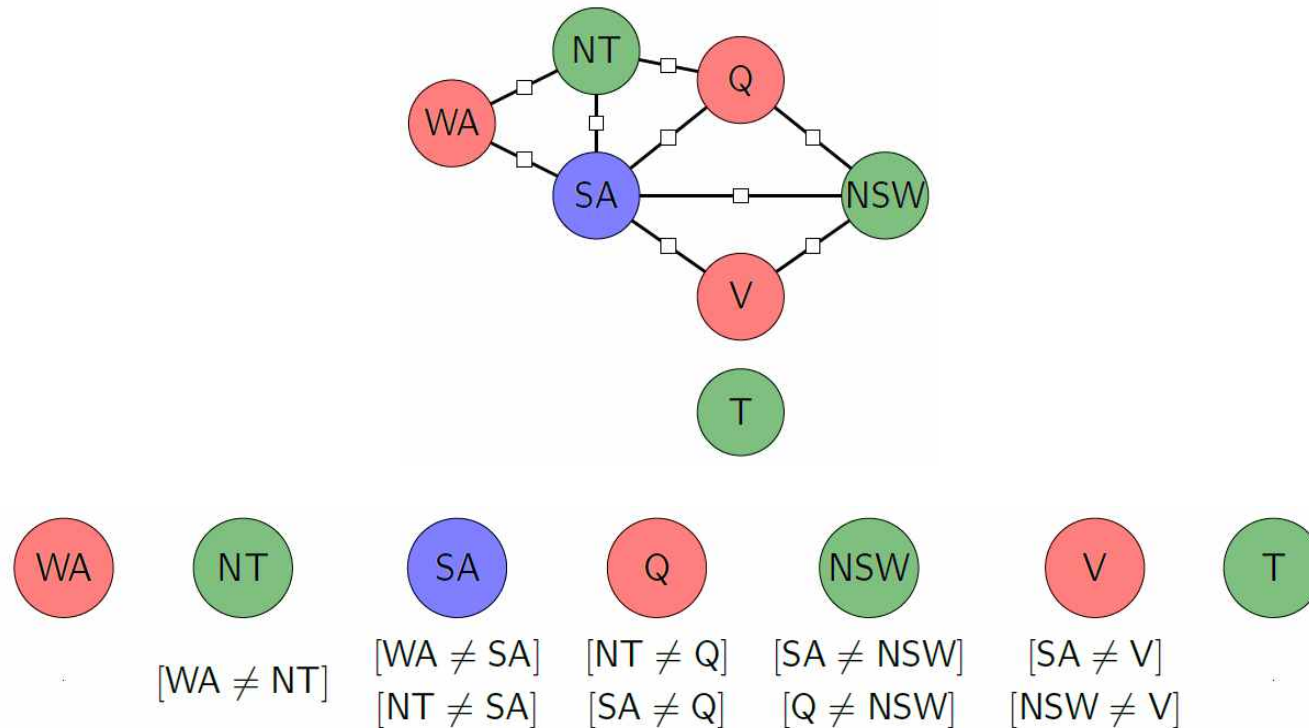
$$\text{Weight}(x) = \prod_{j=1}^m f_j(x)$$

Objective: find the maximum weight assignment

$$\arg \max_x \text{Weight}(x)$$




Partial assignment weights



- Compute weight of partial assignments as we go. (Weight of partial assignment is product of all factors whose scope includes only assigned variables.)

Dependent factors

- Partial assignment (e.g., $x = \{\text{WA: R}, \text{NT: G}\}$)



Definition: dependent factors

- Let $D(x, X_i)$ be set of factors depending on X_i but not on unassigned variables.
- $D(\{\text{WA: R}, \text{NT: G}\}, \text{SA}) = \{[\text{WA} \neq \text{SA}], [\text{NT} \neq \text{SA}]\}$

Backtracking search

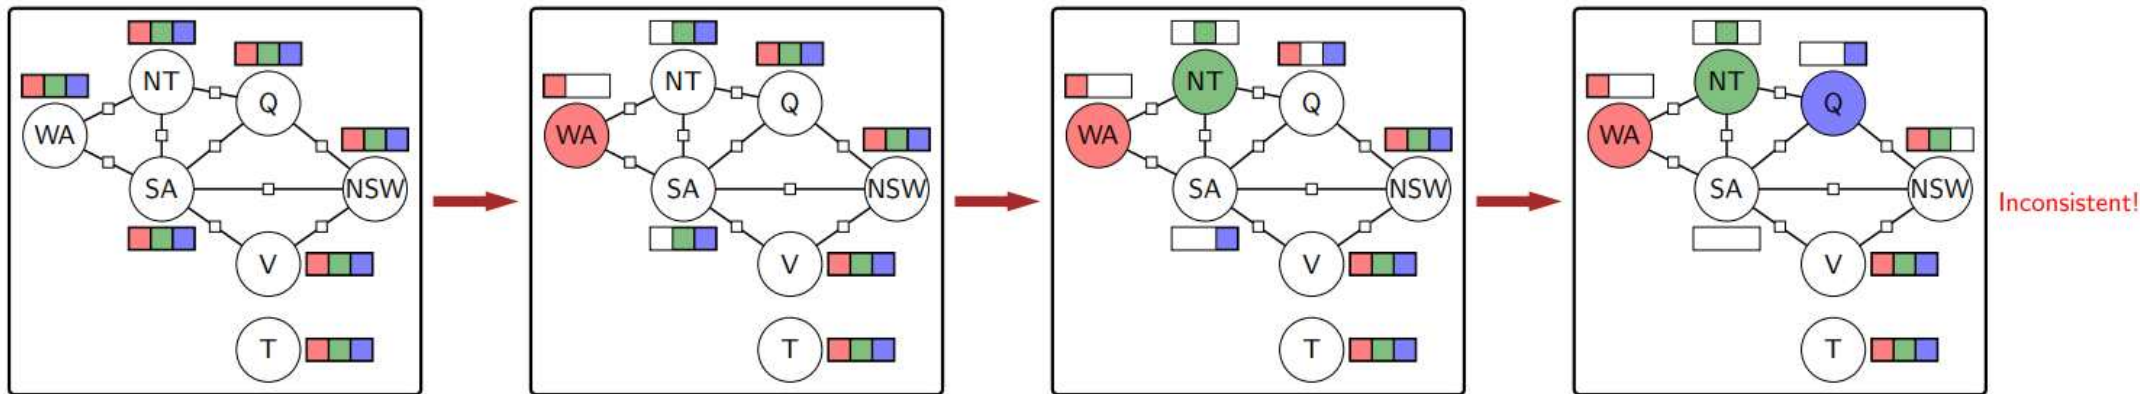
Algorithm: backtracking search

- Backtrack($x, w, \text{Domains}$):
 - If x is complete assignment: update best and return
 - Choose unassigned VARIABLE X_i
 - Order VALUES Domain_i of chosen X_i
 - For each value v in that order:
 - $\delta \leftarrow \prod_{f_j \in D(x, X_i)} f_j(x \cup \{X_i: v\})$
 - If $\delta = 0$: continue
 - $\text{Domains}' \leftarrow \text{Domains via LOOKAHEAD}$
 - If any $\text{Domains}'_i$ is empty: continue
 - Backtrack($x \cup \{X_i: v\}, w\delta, \text{Domains}'$)

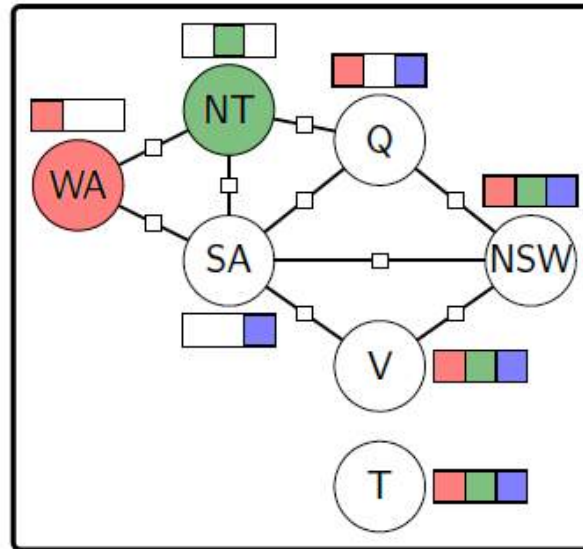
Backtracking search: Lookahead (forward checking)

Key idea: forward checking (one-step lookahead)

- After assigning a variable X_i , eliminate inconsistent values from the domains of X_i 's neighbors.
- If any domain becomes empty, return.



Choosing an unassigned variable



Which variable to assign next?

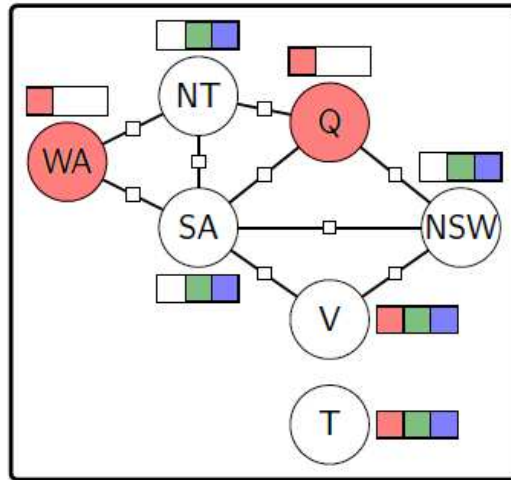
Key idea: most constrained variable

- Choose variable that has the smallest domain.

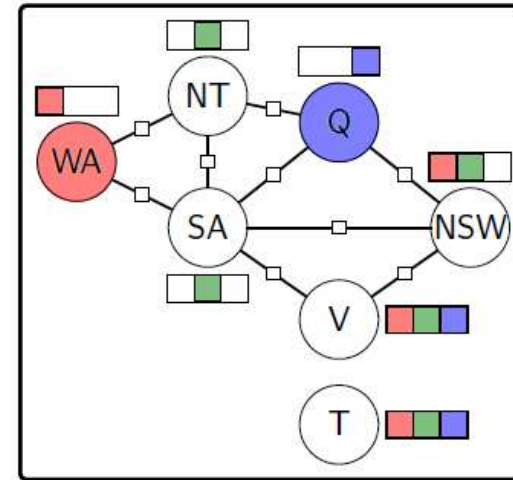
This example: SA (has only one value)

Order values of a selected variable

What values to try for Q?



$2 + 2 + 2 = 6$ consistent values



$1 + 1 + 2 = 4$ consistent values

Key idea: least constrained value

- Order values of selected X_i by decreasing number of consistent values of neighboring variables.

When to fail?

Most constrained **variable** (MCV):

- Must assign **every** variable
- If going to fail, fail early => more pruning

Least constrained **value** (LCV):

- Need to choose **some** value
- Choosing value most likely to lead to solution

When do these heuristics help?

- **Most constrained variable**: useful when **some** factors are constraints (can prune assignments with weight 0)
- **Least constrained value**: useful when **all** factors are constraints (all assignment weights are 1 or 0)
- Forward checking: needed to prune domains to make heuristics useful!

Review: backtracking search with heuristics

Algorithm: backtracking search

- Backtrack($x, w, \text{Domains}$):
 - If x is complete assignment: update best and return
 - Choose unassigned **VARIABLE** X_i (MCV)
 - Order **VALUES** Domain_i of chosen X_i (LCV)
 - For each value v in that order:
 - $\delta \leftarrow \prod_{f_j \in D(x, X_i)} f_j(x \cup \{X_i: v\})$
 - If $\delta = 0$: continue
 - **$\text{Domains}' \leftarrow \text{Domains}$ via LOOKAHEAD** (forward checking)
 - If any $\text{Domains}'_i$ is empty: continue
 - Backtrack($x \cup \{X_i: v\}, w\delta, \text{Domains}'$)

Roadmap

Modeling

Definitions

Examples

Backtracking (exact) search

Dynamic ordering

Arc consistency

Approximate search

Beam search

Local search

Review: backtracking search with heuristics

Algorithm: backtracking search

- Backtrack($x, w, \text{Domains}$):
 - If x is complete assignment: update best and return
 - Choose unassigned **VARIABLE** X_i (MCV)
 - Order **VALUES** Domain_i of chosen X_i (LCV)
 - For each value v in that order:
 - $\delta \leftarrow \prod_{f_j \in D(x, X_i)} f_j(x \cup \{X_i: v\})$
 - If $\delta = 0$: continue
 - **$\text{Domains}' \leftarrow \text{Domains}$ via LOOKAHEAD (AC-3)**
 - If any $\text{Domains}'_i$ is empty: continue
 - Backtrack($x \cup \{X_i: v\}, w\delta, \text{Domains}'$)

Arc consistency

Idea: eliminate values from domains => reduce branching

Example: numbers

- Before enforcing arc consistency on X_i :
 - $X_i \in \text{Domain}_i = \{1,2,3,4,5\}$
 - $X_j \in \text{Domain}_j = \{1,2\}$
 - Factor: $[X_i + X_j = 4]$
- After enforcing arc consistency on X_i :
 - $X_i \in \text{Domain}_i = \{2,3\}$

Arc consistency

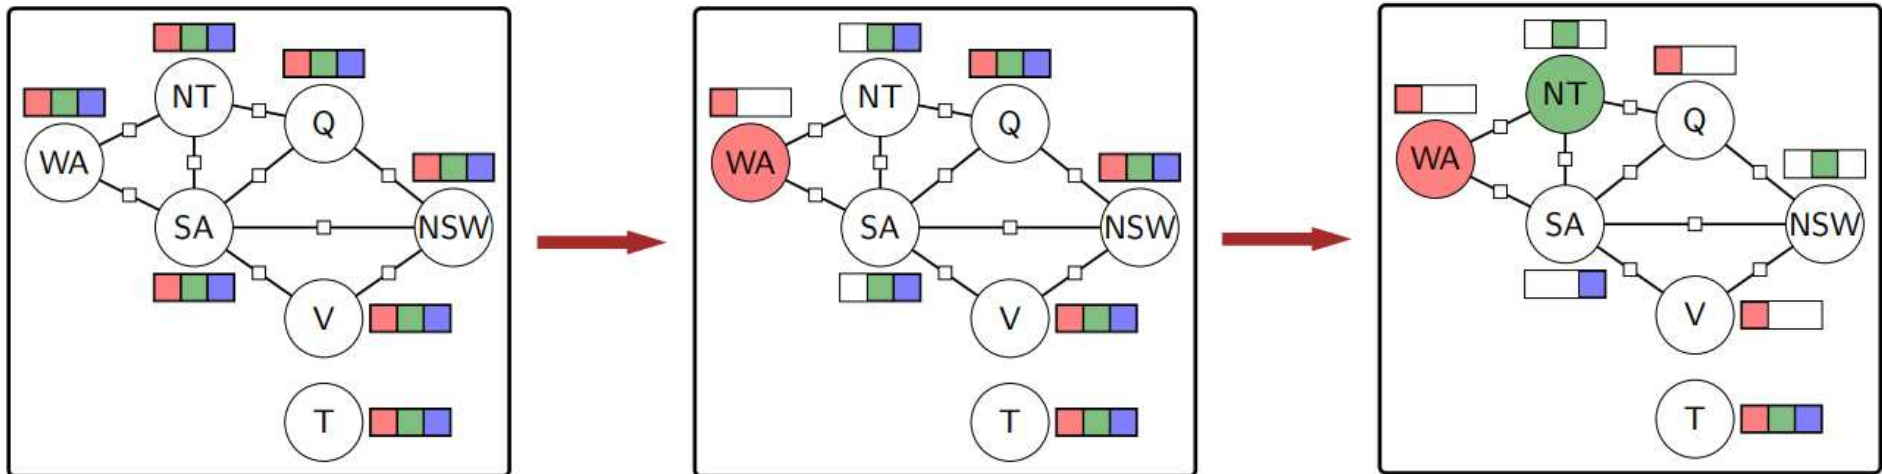
Definition: arc consistency

- A variable X_i is **arc consistent** with respect to X_j if for each $x_i \in \text{Domain}_i$, there exists $x_j \in \text{Domain}_j$ such that $f(\{X_i: x_i, X_j: x_j\}) \neq 0$ for all factors f whose scope contains X_i and X_j .

Algorithm: enforce arc consistency

- $\text{EnforceArcConsistency}(X_i, X_j)$: Remove values from Domain_i to make X_i arc consistent with respect to X_j .

AC-3 (example)



AC-3

Forward checking: when assign $X_j: x_j$, set $\text{Domain}_j = \{x_j\}$ and enforce arc consistency on all neighbors X_i with respect to X_j


AC-3: repeatedly enforce arc consistency on all variables

Algorithm: AC-3

- $S \leftarrow \{X_j\}$.
- While S is non-empty:
 - Remove any X_j from S .
 - For all neighbors X_i of X_j :
 - Enforce arc consistency on X_i w.r.t. X_j .
 - If Domain_i changed, add X_i to S .

Limitations of AC-3

- AC-3 isn't always effective:



- No consistent assignments, but AC-3 doesn't detect a problem!
- Intuition: if we look locally at the graph, nothing blatantly wrong...

Summary

- Enforcing arc consistency: make domains consistent with factors
- Forward checking: enforces arc consistency on neighbors
- AC-3: enforces arc consistency on neighbors and their neighbors, etc.
- Lookahead can speed up backtracking search!

Roadmap

Modeling

Definitions

Examples

Backtracking (exact) search

Dynamic ordering


Arc consistency

Approximate search

Beam search

Local search

Review: factor graph and CSPs



Definition: factor graph

- Variables:

$X = (X_1, \dots, X_n)$, where $X_i \in \text{Domain}_i$

- Factors:

f_1, \dots, f_m , with each $f_j(X) \geq 0$

Definition: assignment weight



- Each **assignment** $x = (x_1, \dots, x_n)$ has a **weight**:

$$\text{Weight}(x) = \prod_{j=1}^m f_j(x)$$

Objective: find the maximum weight assignment

$$\arg \max_x \text{Weight}(x)$$

Example: object tracking



x_1	$o_1(x_1)$
0	2
1	1
2	0

x_2	$o_2(x_2)$
0	0
1	1
2	2

x_3	$o_3(x_3)$
0	0
1	1
2	2

$ x_i - x_{i+1} $	$t_i(x_i, x_{i+1})$
0	2
1	1
2	0

Review: backtracking search

Vanilla version:

$O(|\text{Domain}|^n)$ time

Lookahead: forward checking, AC-3


$O(|\text{Domain}|^n)$ time

Dynamic ordering: most constrained variable, least constrained value


$O(|\text{Domain}|^n)$ time

Note: these pruning techniques useful only for constraints. What if all the factors return strictly positive values?

Backtracking search



Greedy search




Greedy search

Algorithm: greedy search

- Partial assignment $x \leftarrow \{\}$
- For each $i = 1, \dots, n$:
 - Extend:
 - Compute weight of each $x_v = x \cup \{X_i: v\}$
 - Prune:
 - $x \leftarrow x_v$ with highest weight

Not guaranteed to find optimal assignment!

Beam search



Beam size $K = 4$

Beam search

Idea: keep $\leq K$ candidate list C of partial assignments

Algorithm: beam search

- Initialize $C \leftarrow [\{\}]$
- For each $i = 1, \dots, n$:
 - Extend:
 - $C' \leftarrow \{x \cup \{X_i: v\} : x \in C, v \in \text{Domain}_i\}$
 - Prune:
 - $C \leftarrow K$ elements of C' with highest weights

Not guaranteed to find optimal assignment!

Beam search

Time complexity:

- $K = 1$ is greedy ($O(nb)$ time) with branching factor $b = |\text{Domain}|$.
- $K = \infty$ is BFS tree search ($O(b^n)$ time)
- $O(n \cdot (Kb) \cdot \log(Kb))$ with beam size K (since need to sort Kb candidates)
- Beam size K controls tradeoff between efficiency and accuracy
- Backtracking search : DFS :: beam search : pruned BFS

Roadmap

Modeling

Definitions

Examples

Backtracking (exact) search

Dynamic ordering


Arc consistency

Approximate search

Beam search

Local search

Review: factor graph and CSPs



Definition: factor graph

- Variables:

$X = (X_1, \dots, X_n)$, where $X_i \in \text{Domain}_i$

- Factors:

f_1, \dots, f_m , with each $f_j(X) \geq 0$

Definition: assignment weight

- Each **assignment** $x = (x_1, \dots, x_n)$ has a **weight**:

$$\text{Weight}(x) = \prod_{j=1}^m f_j(x)$$

Objective: find the maximum weight assignment

$$\arg \max_x \text{Weight}(x)$$

Local search


- Backtracking/beam search: extend partial assignments




- Local search: modify complete assignments



Example: object tracking



One small step



Current assignment: (0, 0, 1); how to improve?

(x_1, \mathbf{v}, x_3) weight


$$(0, \mathbf{0}, 1) \quad 2 \cdot \mathbf{2} \cdot \mathbf{0} \cdot \mathbf{1} \cdot 1 = 0$$

$$(0, \mathbf{1}, 1) \quad 2 \cdot \mathbf{1} \cdot \mathbf{1} \cdot \mathbf{2} \cdot 1 = 4$$

$$(0, \mathbf{2}, 1) \quad 2 \cdot \mathbf{0} \cdot \mathbf{1} \cdot \mathbf{1} \cdot 1 = 0$$

New assignment: (0, 1, 1)

Exploiting locality



Weight of new assignment (x_1, v, x_3)

$$o_1(x_1) t_1(x_1, v) o_2(v) t_2(v, x_3) o_3(x_3)$$




Key idea: locality

- When evaluating possible re-assignments to X_i , only need to consider the factors that depend on X_i .

Iterated conditional modes (ICM)


Algorithm: iterated conditional modes (ICM)

- Initialize x to a random complete assignment
- Loop through $i = 1, \dots, n$ until convergence:
 - Compute weight of $x_v = x \cup \{X_i: v\}$ for each v
 - $x \leftarrow x_v$ with highest weight




Convergence properties

- $\text{Weight}(x)$ increases or stays the same each iteration
- Converges in a finite number of iterations
- Can get stuck in **local optima**
- Not guaranteed to find optimal assignment!



Summary



Algorithms for max-weight assignments in factor graphs:

(1) Extend partial assignments:

- Backtracking search: exact, exponential time
- Beam search: approximate, linear time

(2) Modify complete assignments:

- Iterated conditional modes: approximate, linear time, deterministic
- (Markov networks) Gibbs sampling: approximate, linear time, randomized